QUERY OPTIMIZATION

Vaibhav Bajpai

Contents

Overview

* Optimization
  * find a better equivalent plan
  * ideally want to find a best plan
  * practically want to avoid worst plans

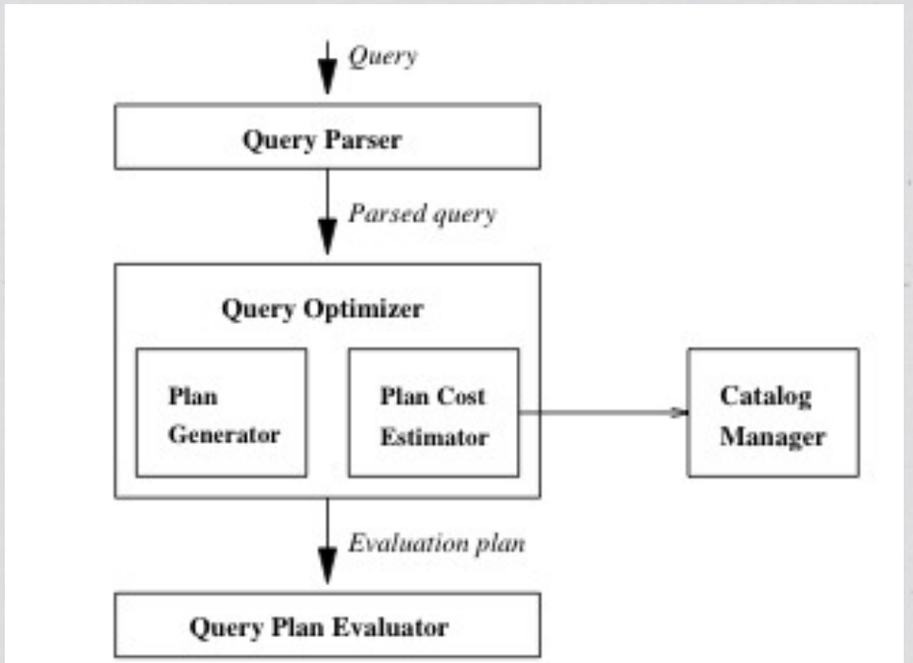

Overview

* Blocking Operator
  * Materialization
* Non-Blocking Operator
  * Pipelining

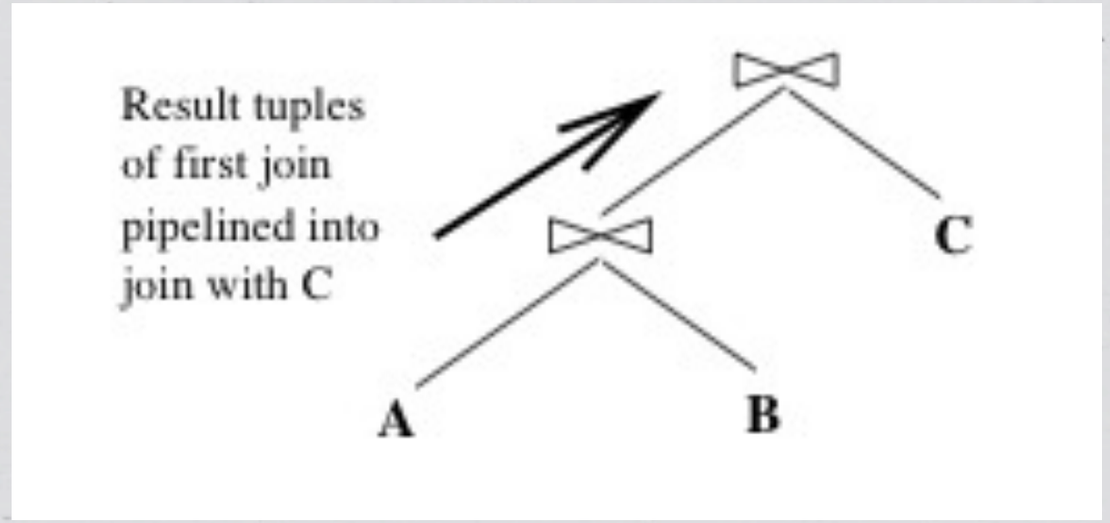

Heuristic-based Approach

(Logical-Level Optimization)

* applying algebraic-equivalences …
* though cannot determine the **actual** costs …
* allows to **push** selection and projection ahead of joins …

Heuristic-based Approach

Equivalences

✻ Selection

$$\pi_a(\sigma_c(R)) \equiv \sigma_c(\pi_a(R))$$

$$\sigma_c(R \times S) \equiv \sigma_c(R) \times S$$

$$\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S$$

$$\sigma_{c_1}(\sigma_{c_2}(\sigma_{c_3}(R \times S)))$$

$$R \bowtie_c S \equiv \sigma_c(R \times S)$$

$$\sigma_c(R \times S) \equiv \sigma_{c_1 \wedge c_2 \wedge c_3}(R \times S)$$

$$\sigma_{c_1}(\sigma_{c_2}(R) \times \sigma_{c_3}(S)).$$

✻ Projection

$$\pi_a(R \times S) \equiv \pi_{a_1}(R) \times \pi_{a_2}(S)$$

$$\pi_a(R \bowtie_c S) \equiv \pi_{a_1}(R) \bowtie_c \pi_{a_2}(S)$$

Cost-based Approach

(Physical-Level Optimization)

* input: logical-plan + physical situation + histograms
* output: efficient physical plan
* approach: **enumerate** all physical plans and pick the best one

Cost-based Approach

System R Optimizer

* works well with less than **10 joins**
* considers both combination of **CPU and IO** costs
* **cartesian products** are avoided
* **only left deep plans** are considered

Cost-based Approach

Query Blocks: Unit of Optimization

* A query is parsed into collection of query blocks
* A query evaluation plan is chosen for each block!
* A query block contains **no nesting** -
    * exactly **one** SELECT and **one** FROM
    * atmost **one** WHERE, GROUP BY, HAVING

Cost-based Approach

Enumeration of Alternative Plans

* Single Relation Plans
  * Clustered Index — (NPages(X)+NPages(Y)) * RF
  * Unclustered Index — (NPages(X)+NTuples(Y)) * RF
* Multiple Relation Plans
  * No. of plans increase exponentially with increase in joins.
  * Need to **restrict search space! (ONLY left-deep join!)**

Cost-based Approach
Enumeration of Alternative Plans
Left Deep Joins

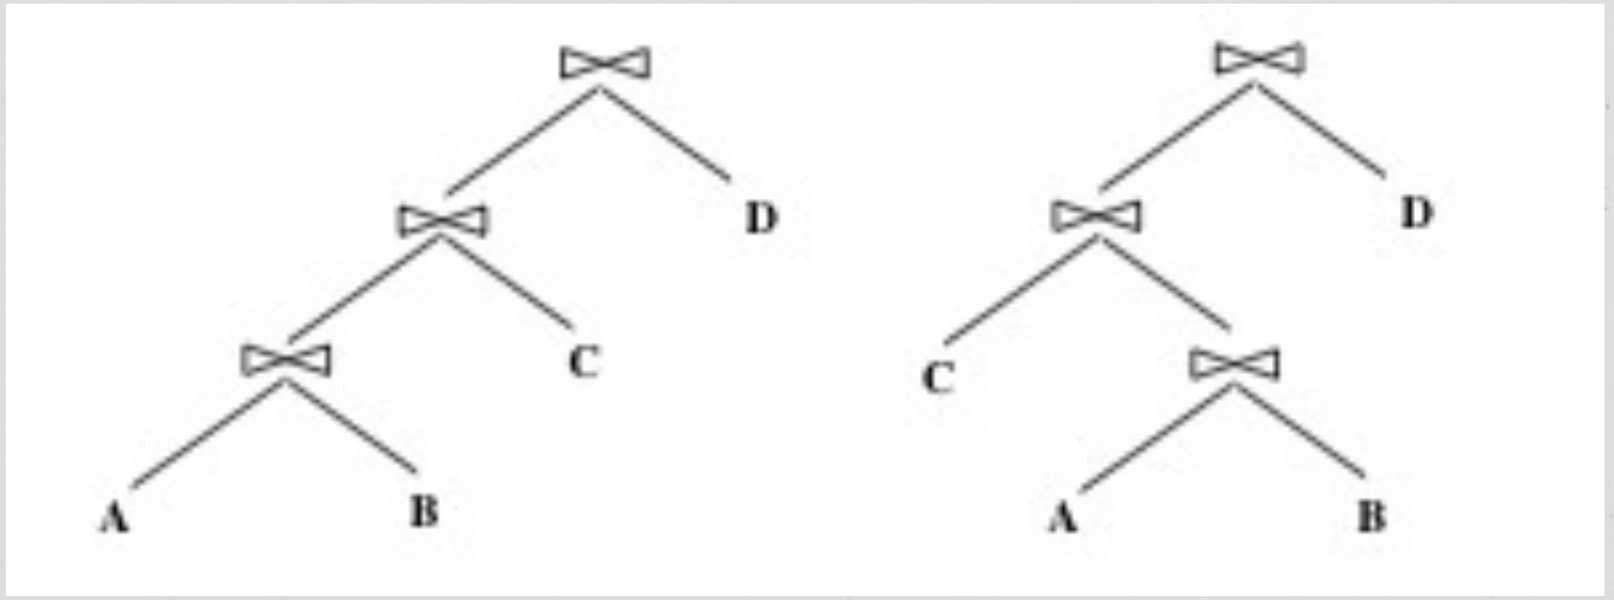

* Why?
  * generates **fully pipelined plans!**
  * many join algorithms scan the right relation multiple times

Example!

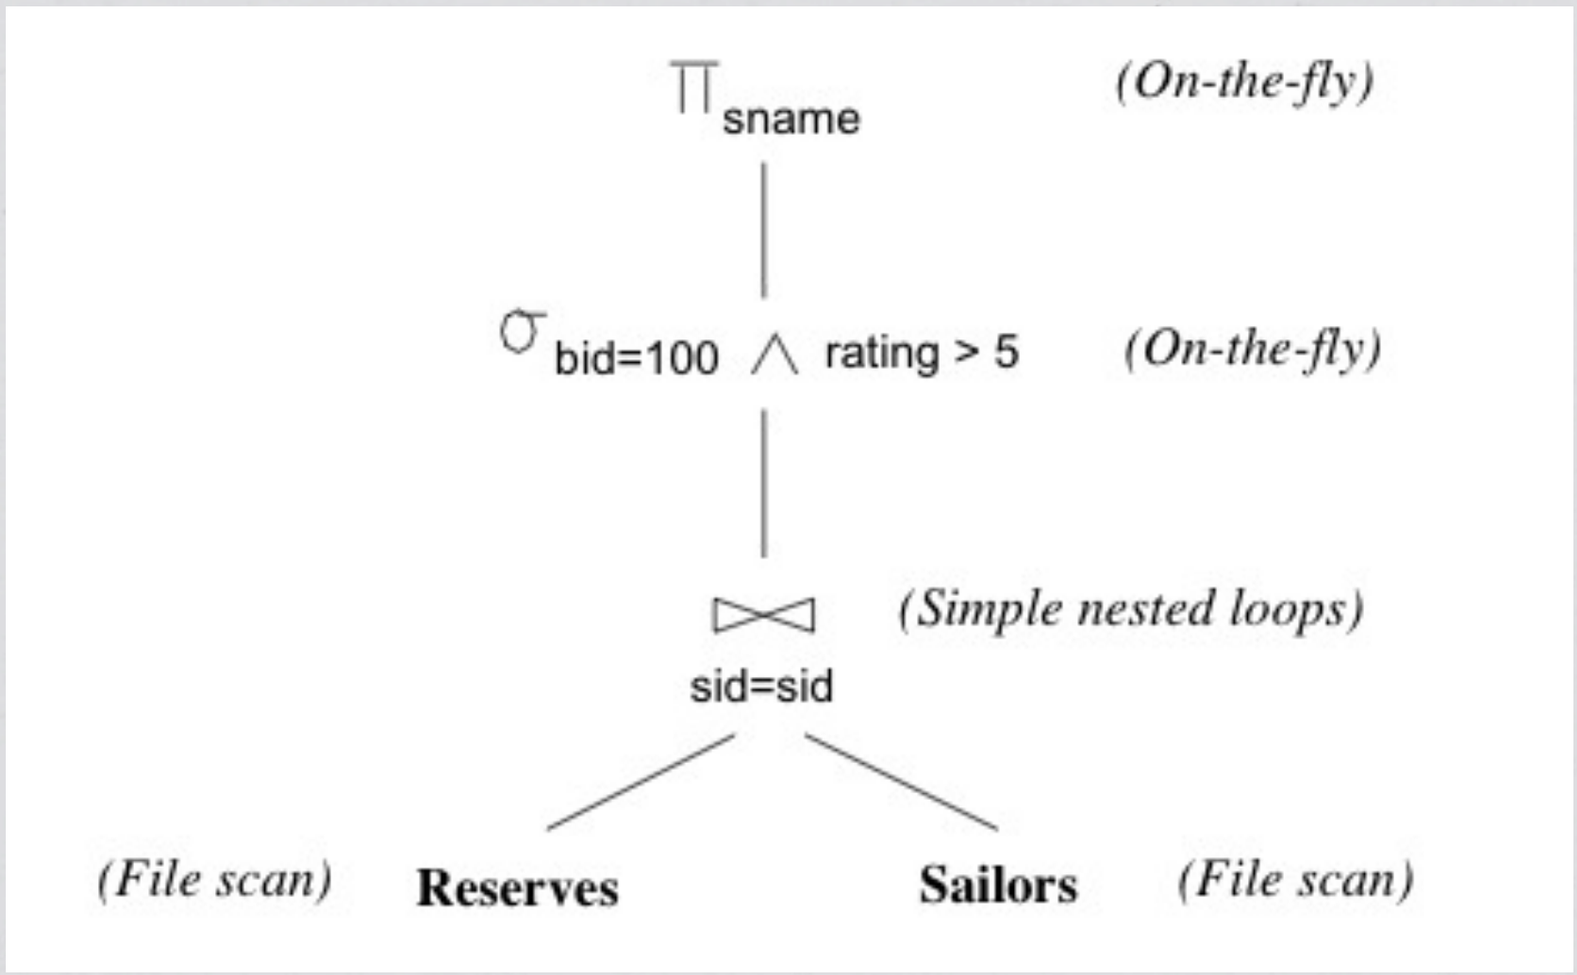
Π sname
(On-the-fly)
σ bid=100 ∧ rating > 5
(On-the-fly)
⋈
(Simple nested loops)
sid=sid
(File scan)
Reserves
Sailors
(File scan)

Example!

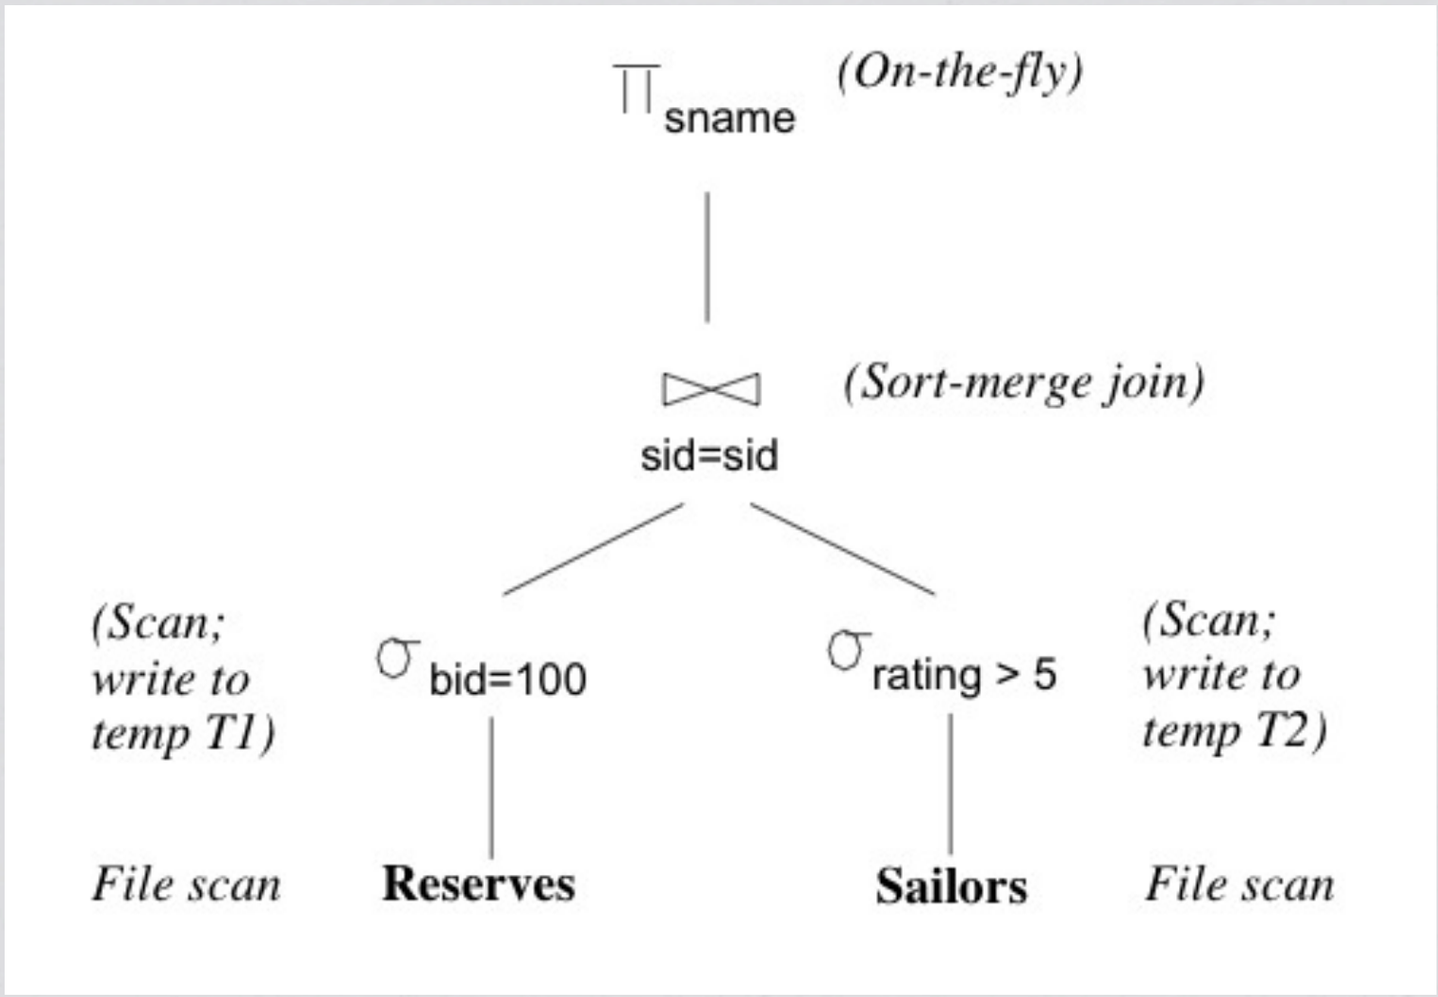
π sname
(On-the-fly)
⋈
(Sort-merge join)
sid=sid
(Scan; write to temp T1)
σ bid=100
σ rating > 5
(Scan; write to temp T2)
File scan
Reserves
Sailors
File scan

Example!

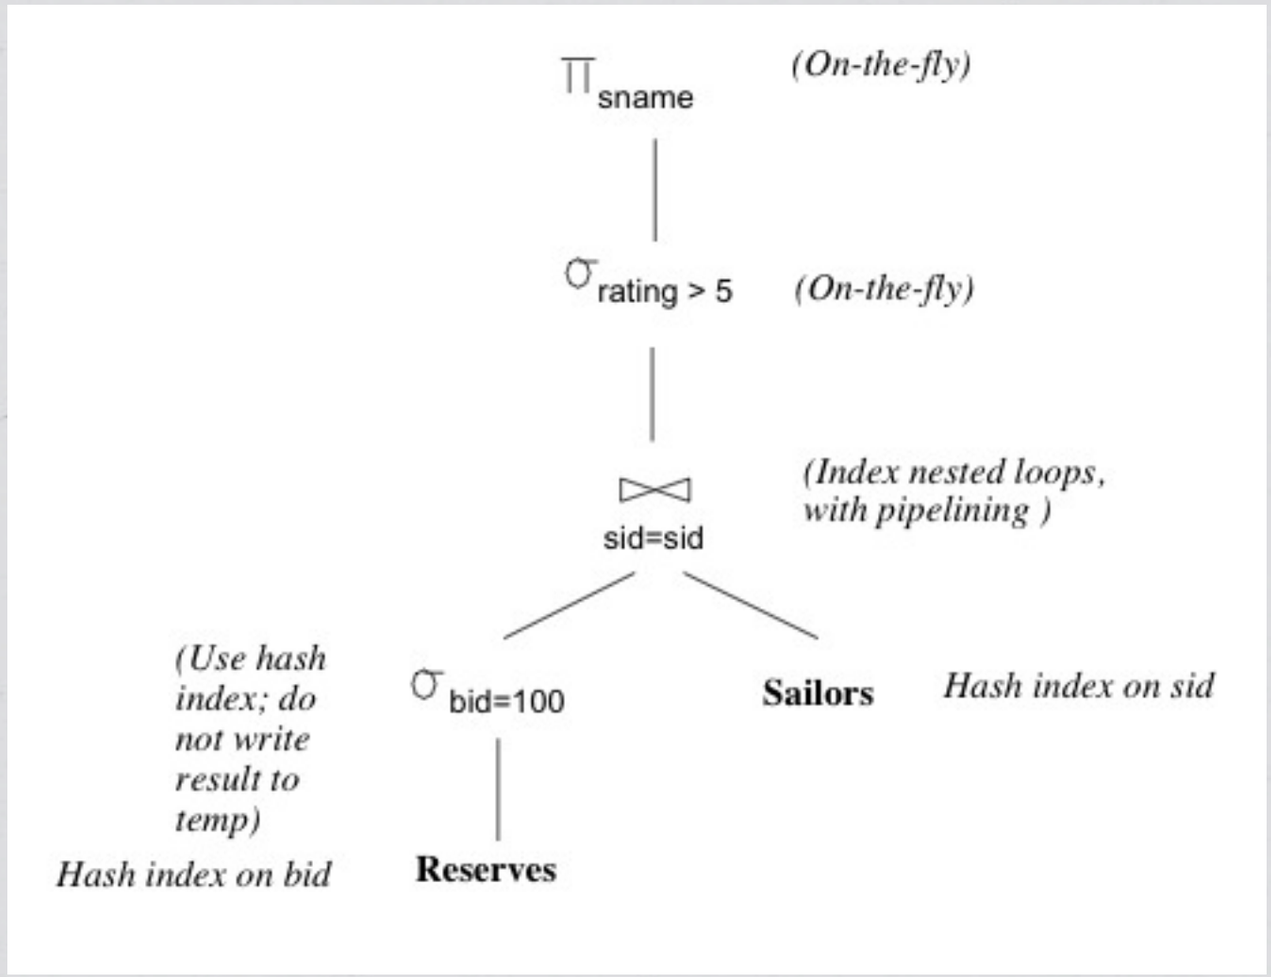
π sname
(On-the-fly)
σ rating > 5
(On-the-fly)
(Index nested loops, with pipelining)
sid=sid
(Use hash index; do not write result to temp)
σ bid=100
Sailors
Hash index on sid
Hash index on bid
Reserves