

# Evaluating the Effectiveness of Happy Eyeballs

Vaibhav Bajpai and Jürgen Schönwälder

RIPE66, Dublin

Computer Networks and Distributed Systems  
Jacobs University Bremen  
Bremen, Germany

May 2013

# overview

- ⌚ motivation
- ⌚ metric and implementation
- ⌚ measurement trials
- ⌚ data analysis insights
- ⌚ conclusion

# motivation

- ⦿ `getaddrinfo(...)` behavior:
  - ⦿ returns list of endpoints in an order that prioritizes IPv6 upgrade path
  - ⦿ order is dictated by [RFC 6724] and `/etc/gai.conf`
  - ⦿ if IPv6 is broken, application is unresponsive in order of seconds
- ⦿ happy eyeballs algorithm [RFC 6555]:
  - ⦿ initiate a TCP `connect(...)` with the first endpoint, give it 300ms
  - ⦿ switch over with a TCP `connect(...)` to a different address family otherwise
  - ⦿ the competition runs fair after 300ms
- ⦿ does the algorithm help improve the user experience?

# metric and implementation

- developed a simple TCP happy eyeballs [RFC 6555] probing tool

```
( service name, port ) —→ happy —→ connection establishment time for each endpoint of a service (μsecs)
```

- uses `getaddrinfo(...)` to resolve service names to endpoints
- uses non-blocking `connect(...)` to connect to all endpoints of a service
- uses a short-delay between connection attempts to avoid SYN floods
- the service name resolution time is not accounted in the output
- can produce either human-readable or machine-readable output
- file locking capability

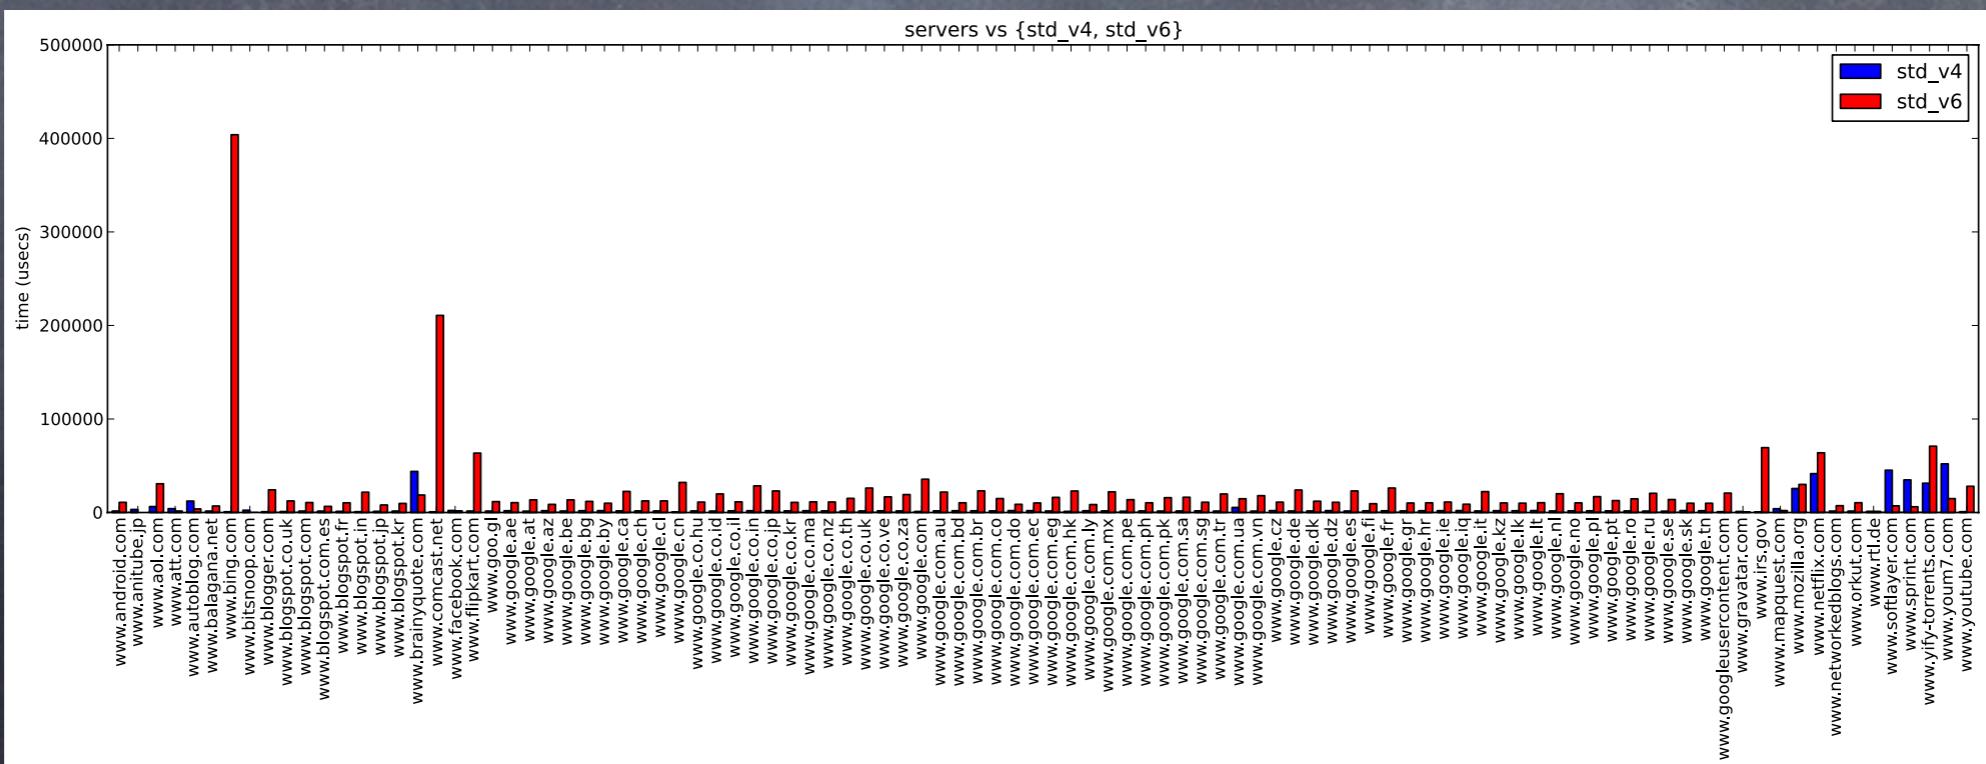
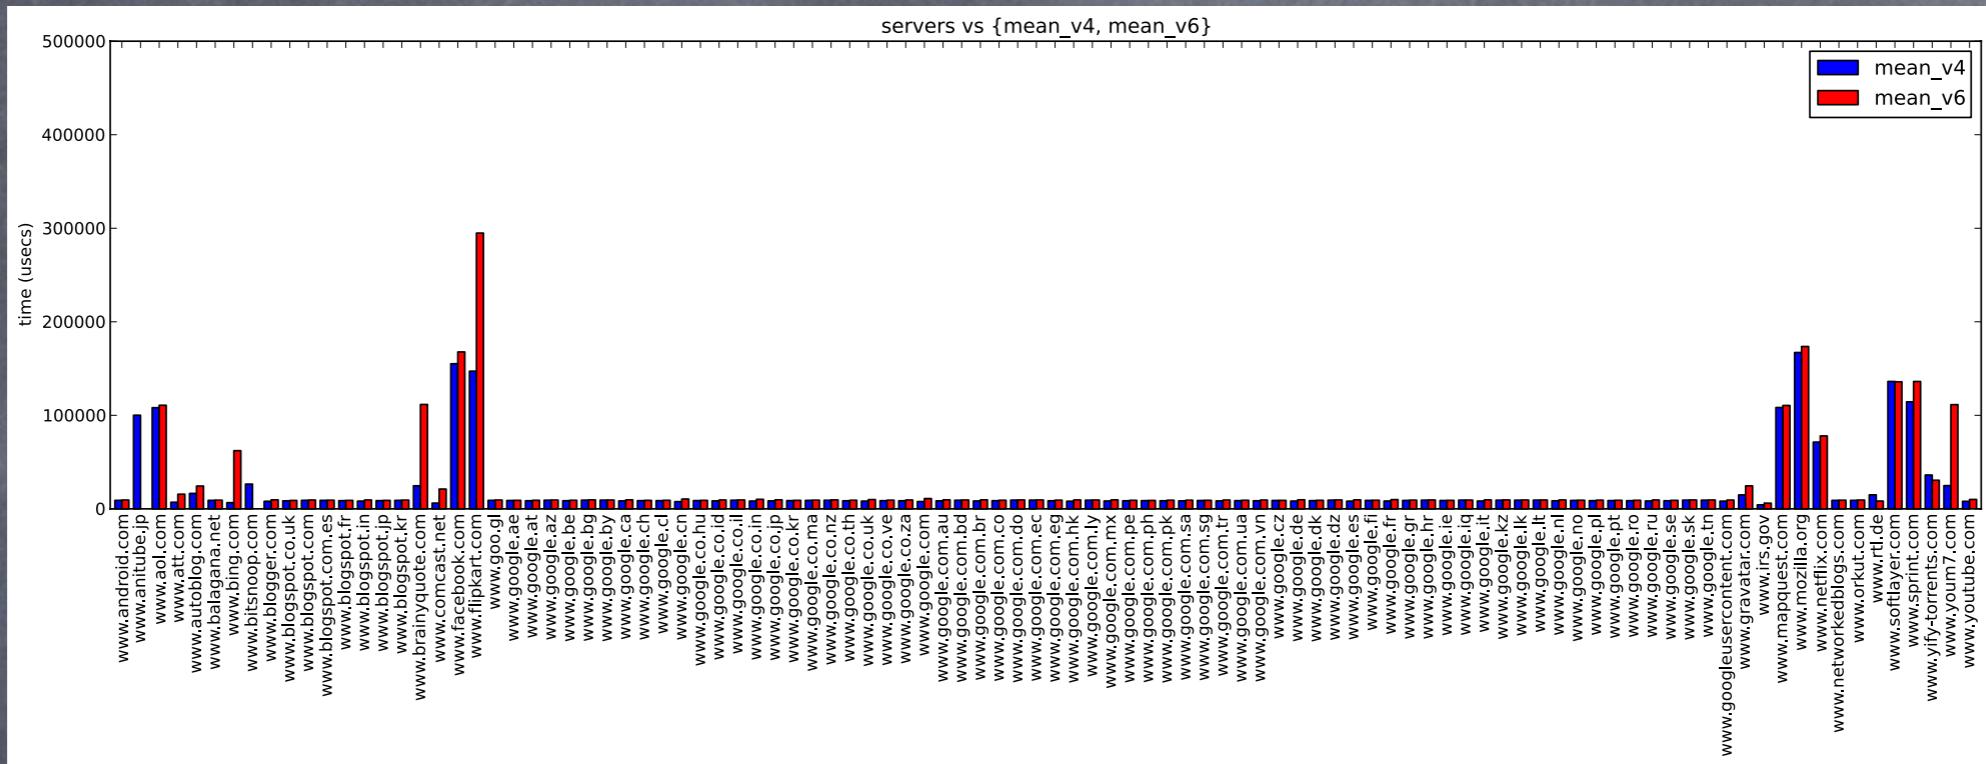
```
>> ./happy -q 1 -m www.google.com www.facebook.com
HAPPY.0;1360681039;0K;www.google.com;80;173.194.69.105;8626
HAPPY.0;1360681039;0K;www.google.com;80;2a00:1450:4008:c01::69;8884
HAPPY.0;1360681039;0K;www.facebook.com;80;2a03:2880:10:6f01:face:b00c::8;170855
HAPPY.0;1360681039;0K;www.facebook.com;80;31.13.72.39;26665
```

# measurement trials

- ⦿ dual-stacked web service name list:
  - ⦿ HE.net maintains a list of top 100 dual-stacked service names
    - ⦿ they use 1M service names from Alexa Top Sites
    - ⦿ some domains we expect are missing from the list
    - ⦿ some services only provide a IPv6 endpoint or prepending a www
    - ⦿ HE.net does not follow CNAMEs (for e.g. wikipedia.org)
  - ⦿ amazon has made 1M service name list public
    - ⦿ we use it and script it ourselves to explicitly follow CNAMEs
- ⦿ measurement agents:
  - ⦿ native IPv6, 6in4, Teredo, IPv6 tunnel broker endpoints, native IPv4
  - ⦿ located at Bremen, Amsterdam, Braunschweig
- ⦿ measurement cycle length:
  - ⦿ 1 week

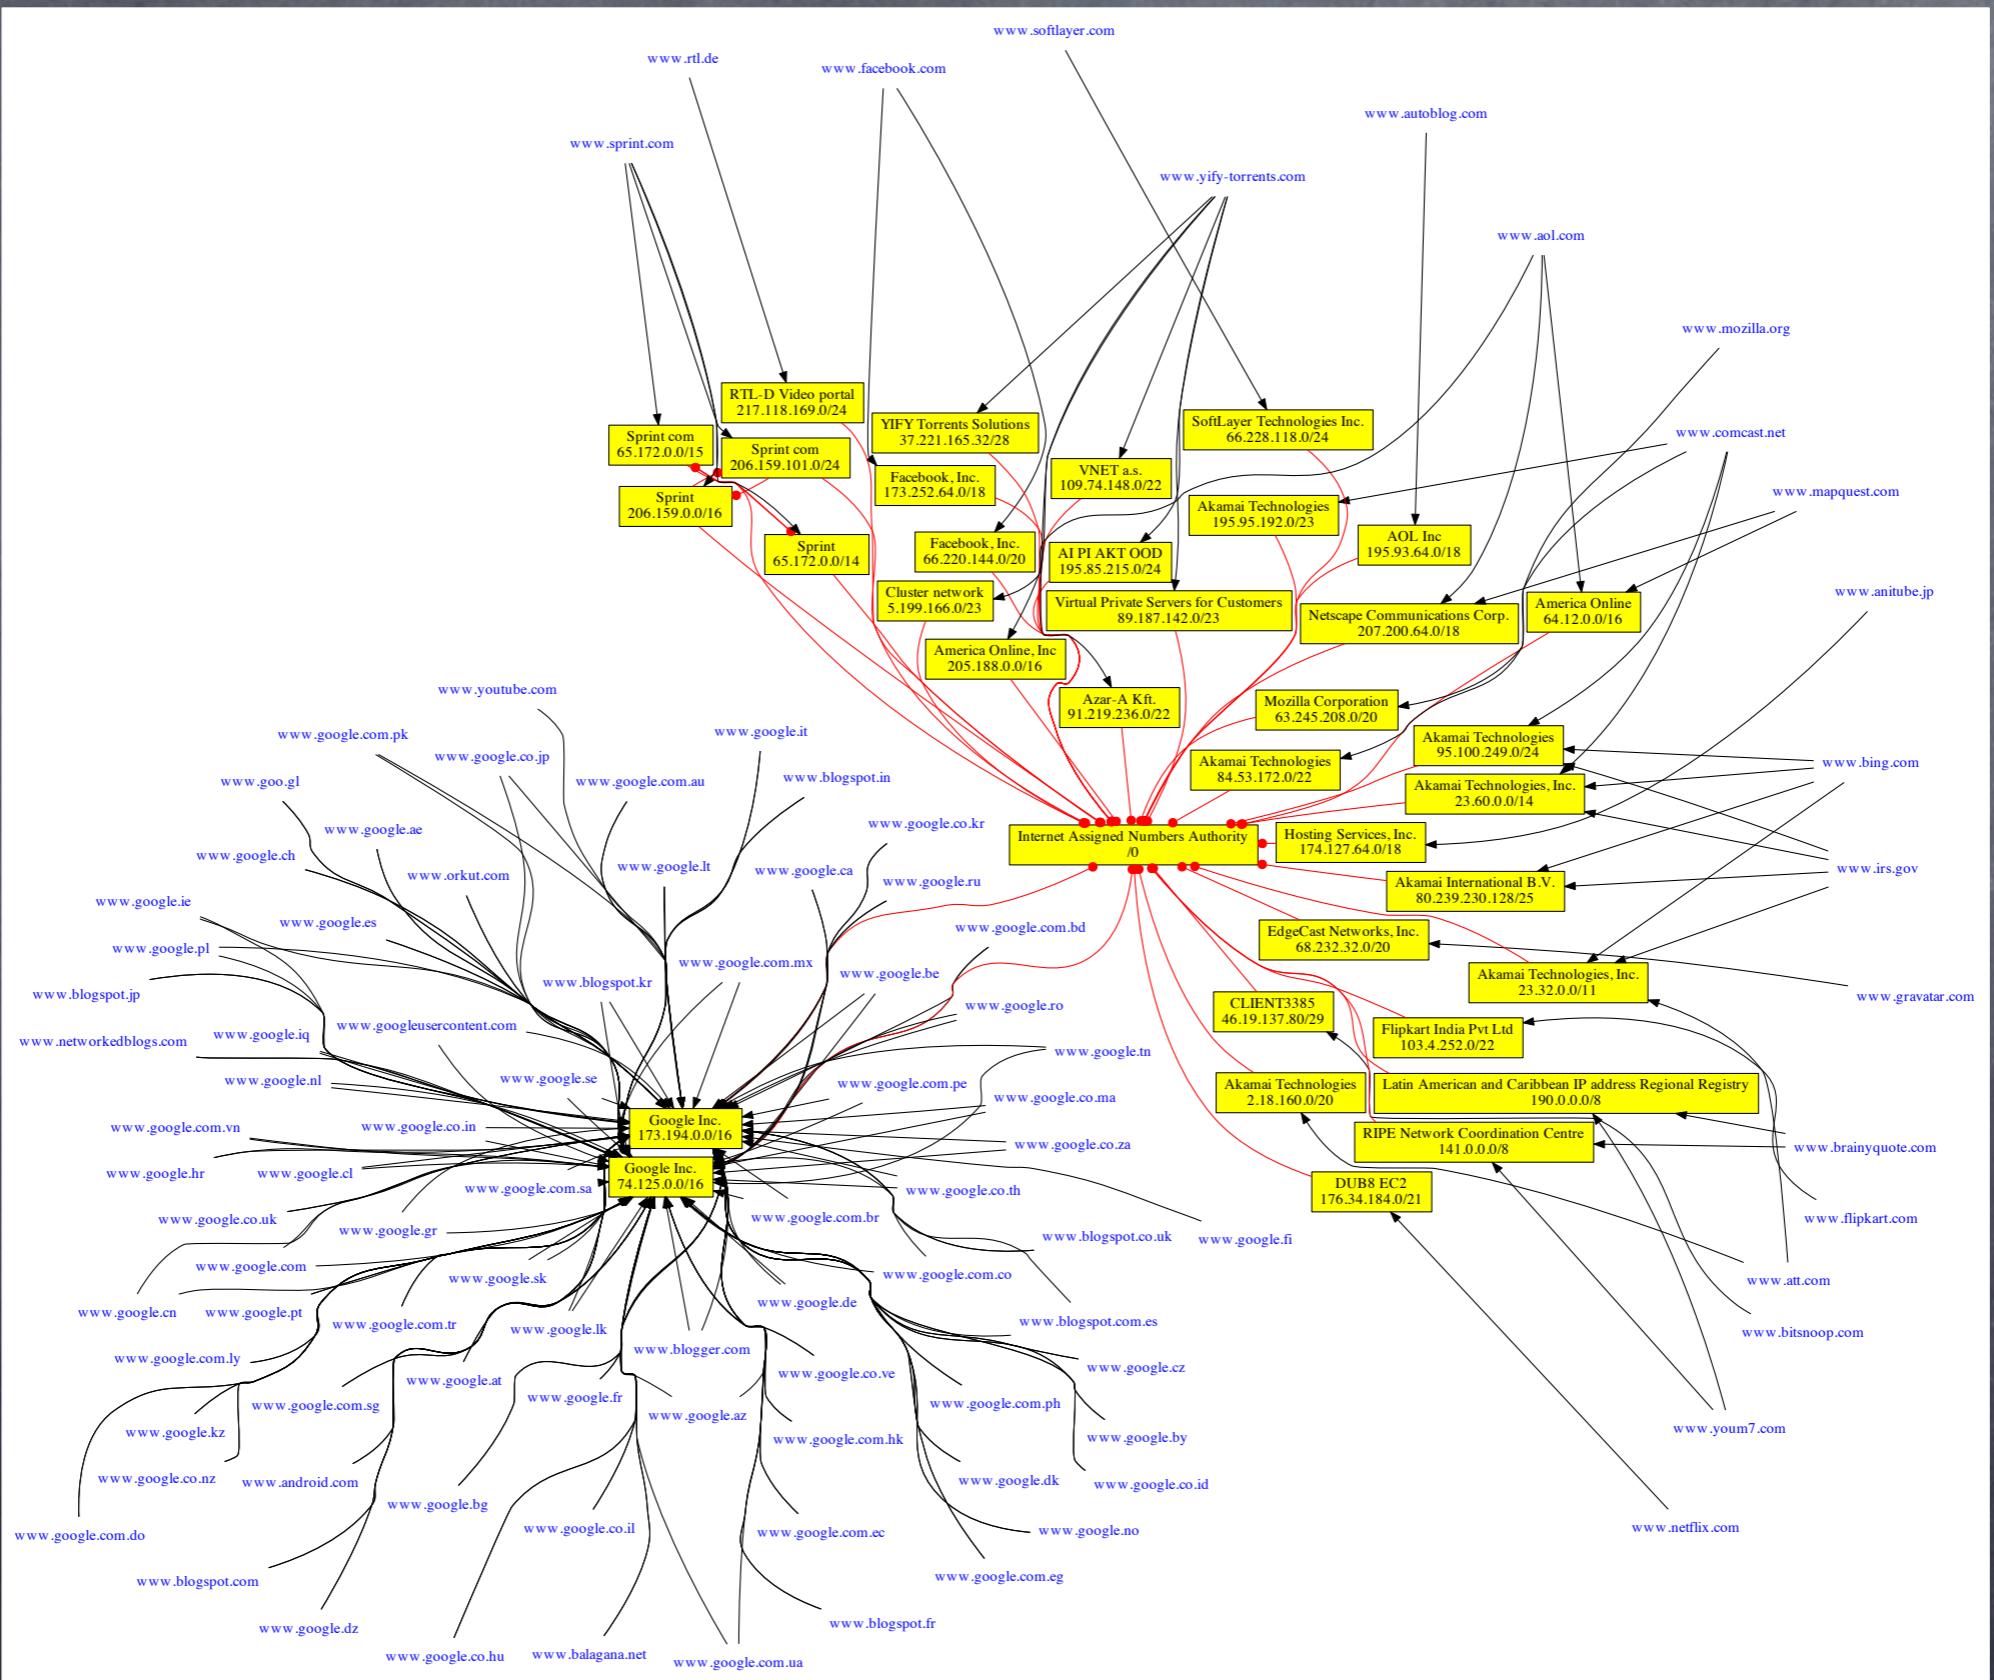
how does IPv6 compare in  
performance to IPv4?

# connection times (mean/std)

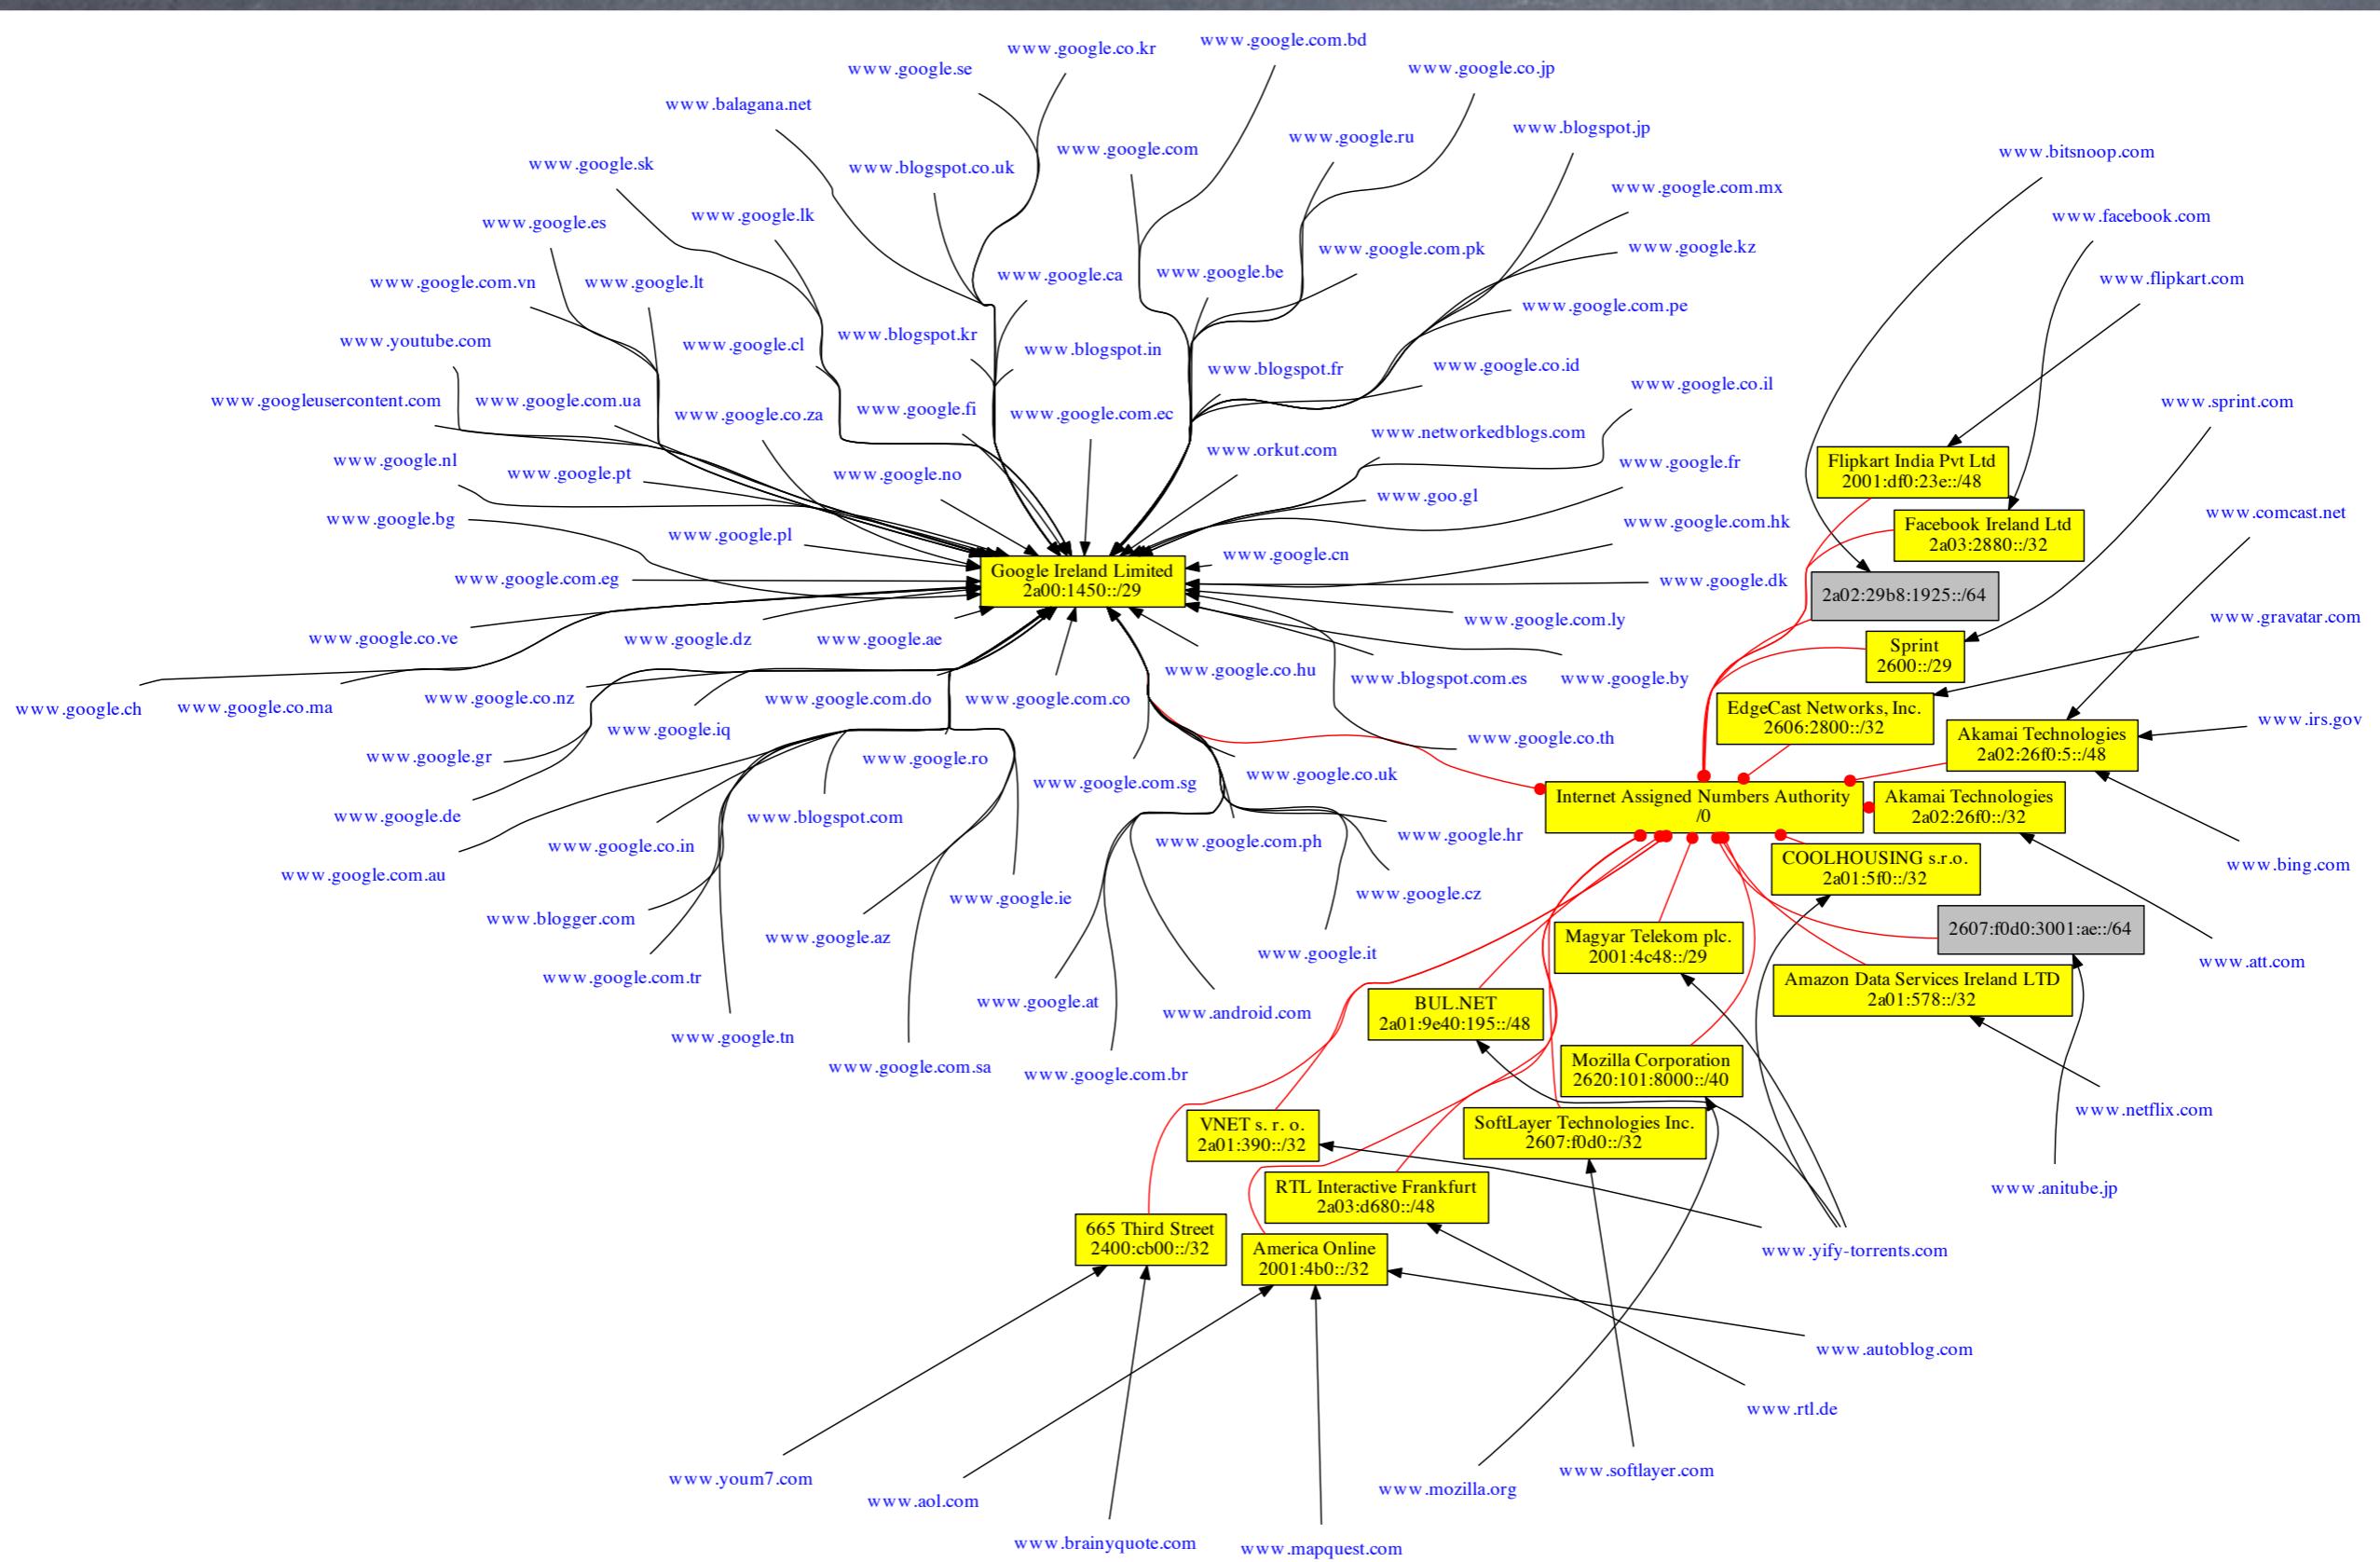


do a major portion of services  
centralize on CDNs?

# IPv4 aggregation cloud



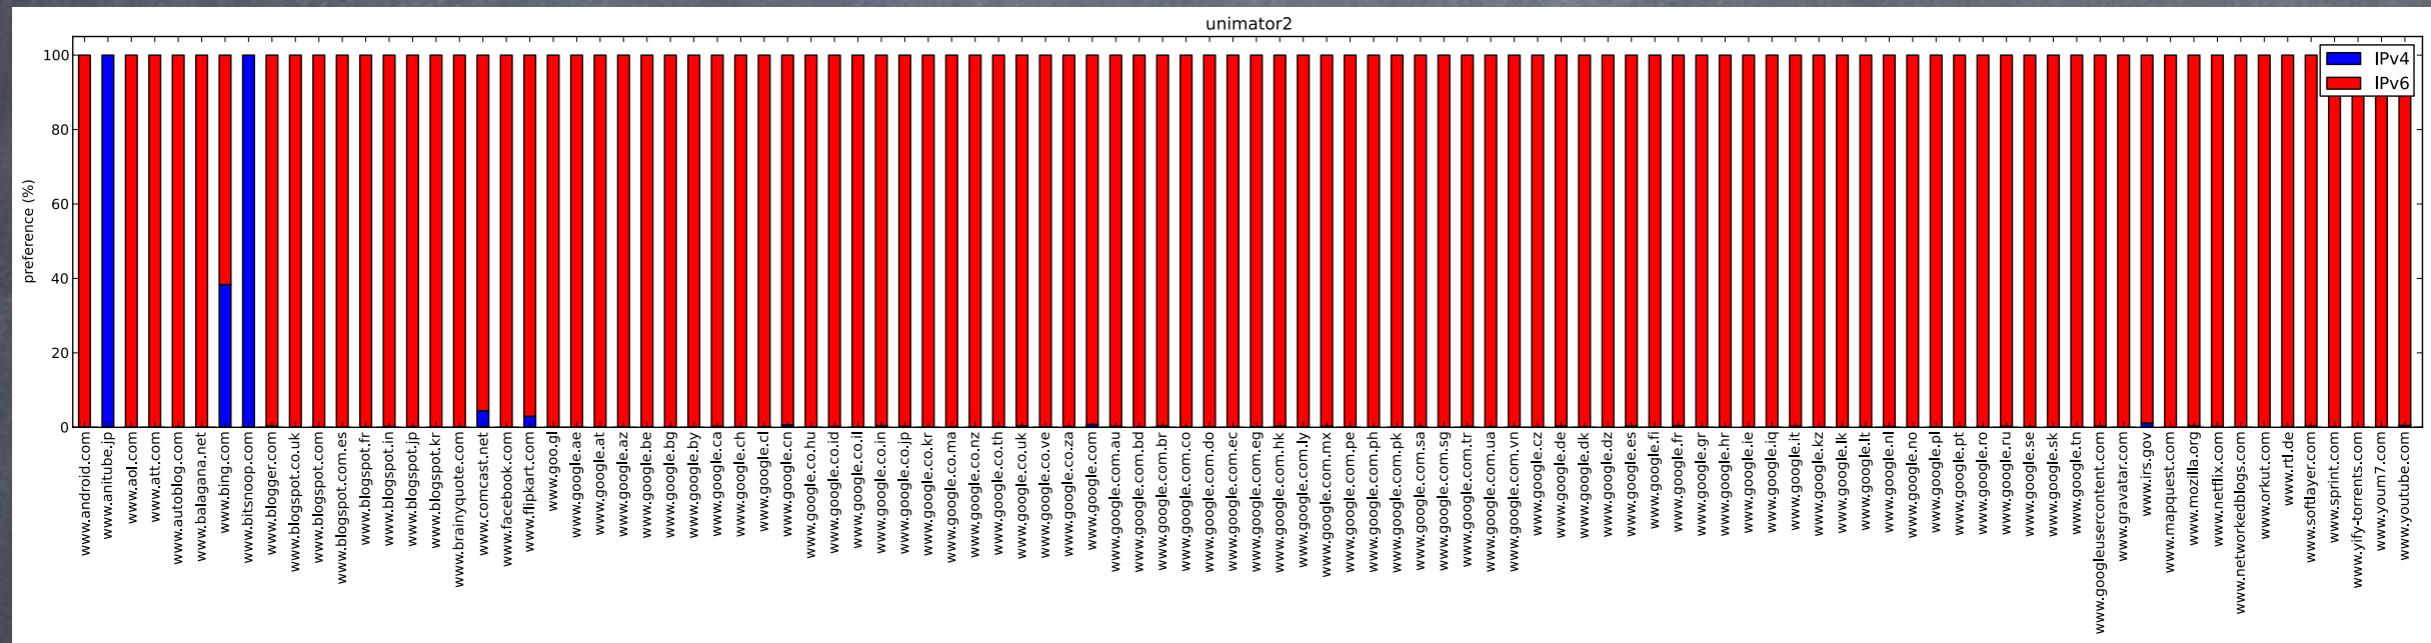
# IPv6 aggregation cloud



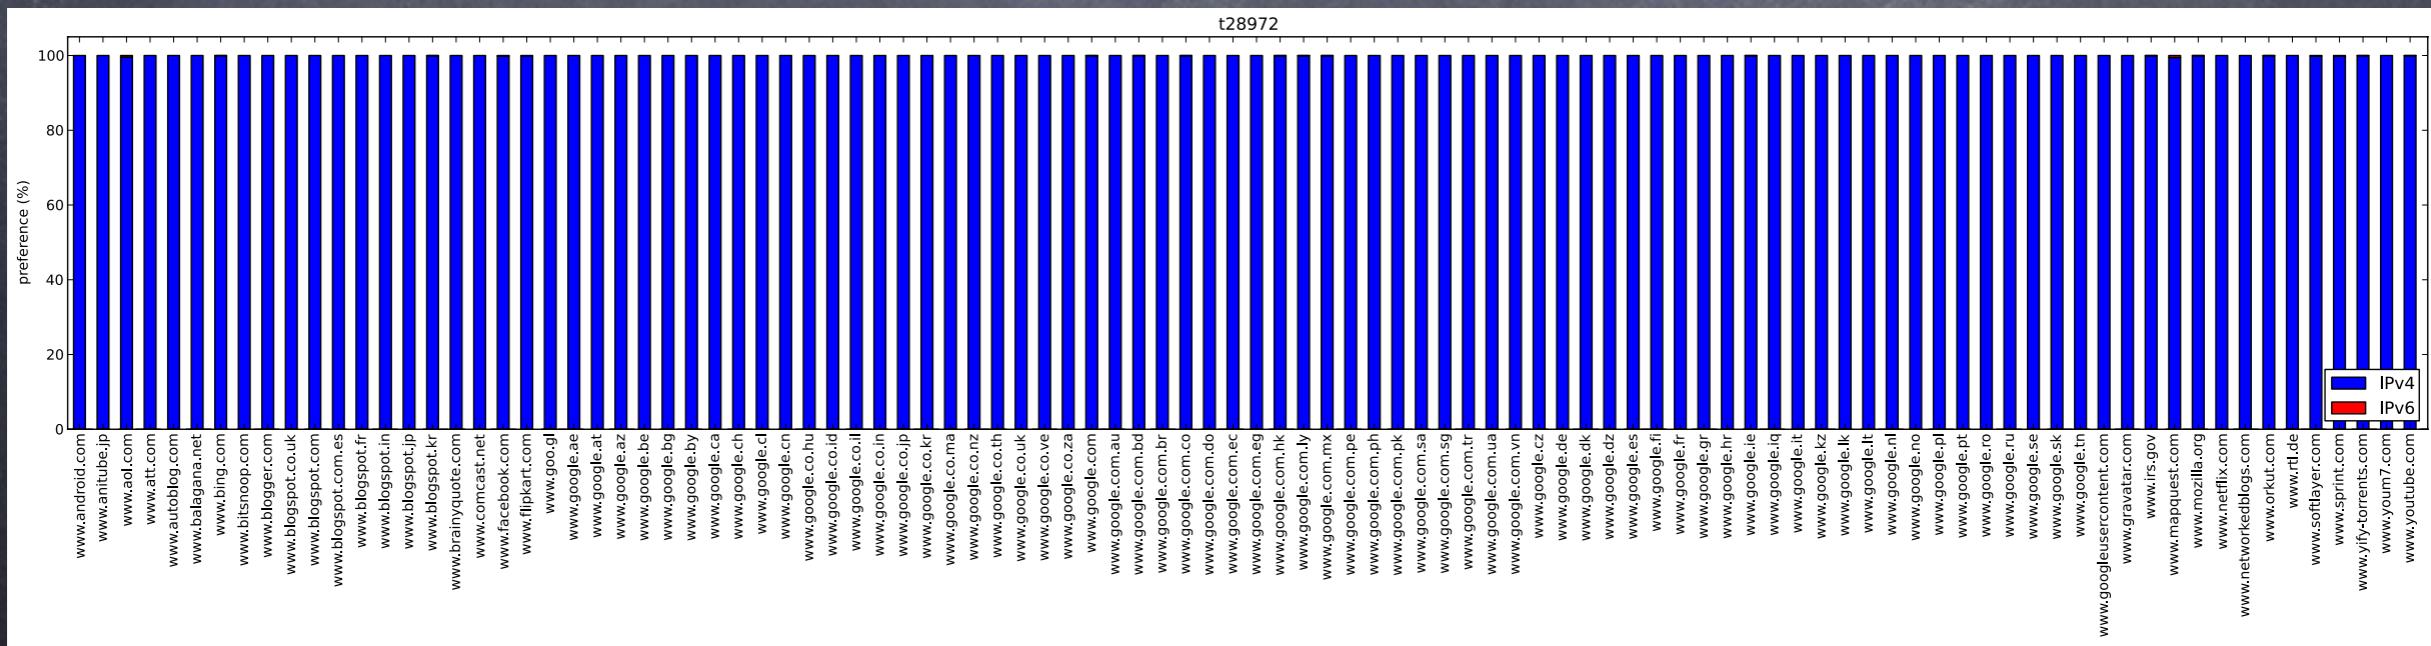
to what extent is IPv6 preferred  
when connecting to a dual-  
stacked service?

# IPv6 preference %

## Native IPv6



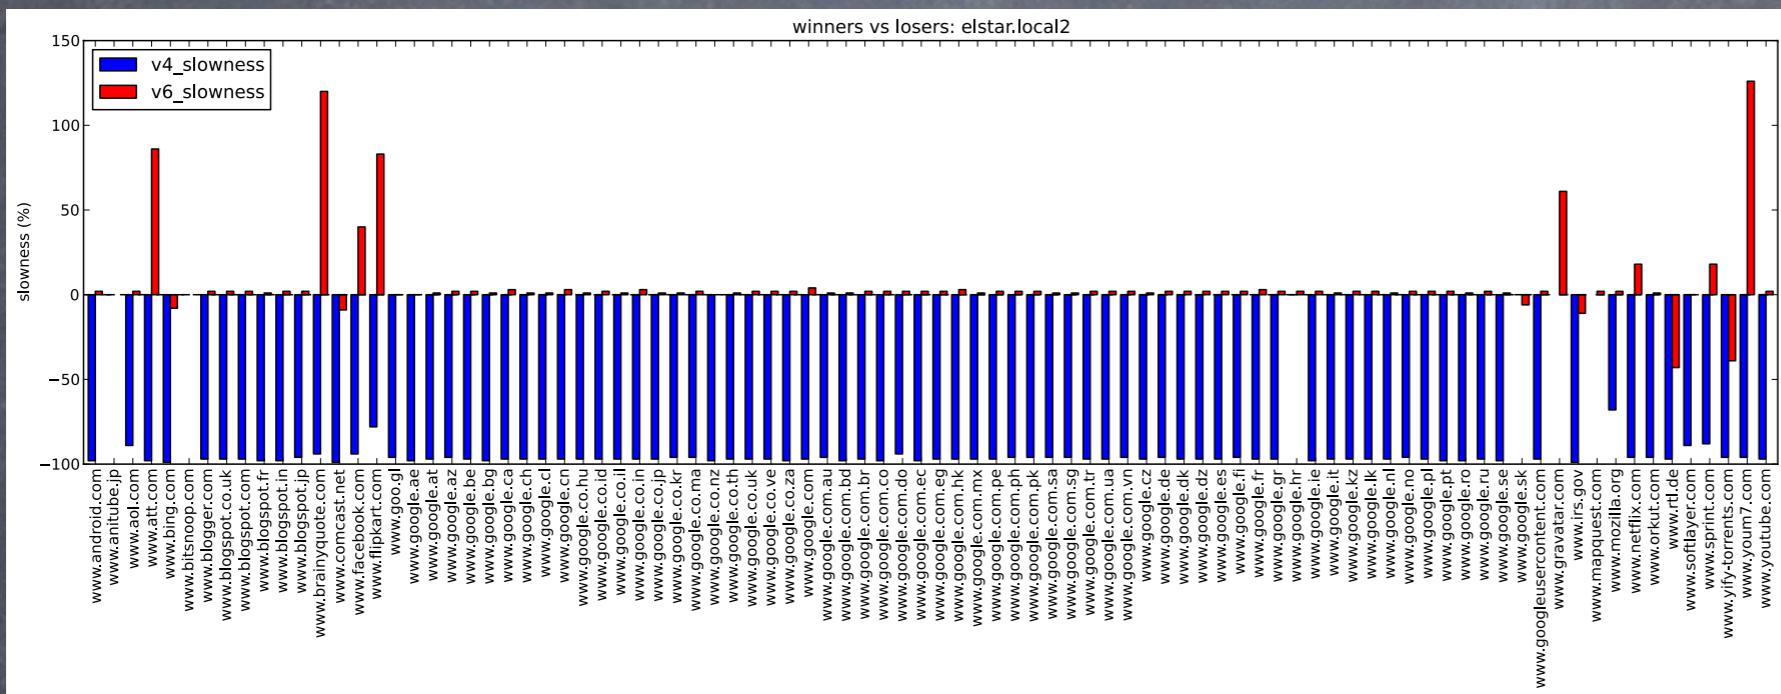
## IPv6 teredo tunnel



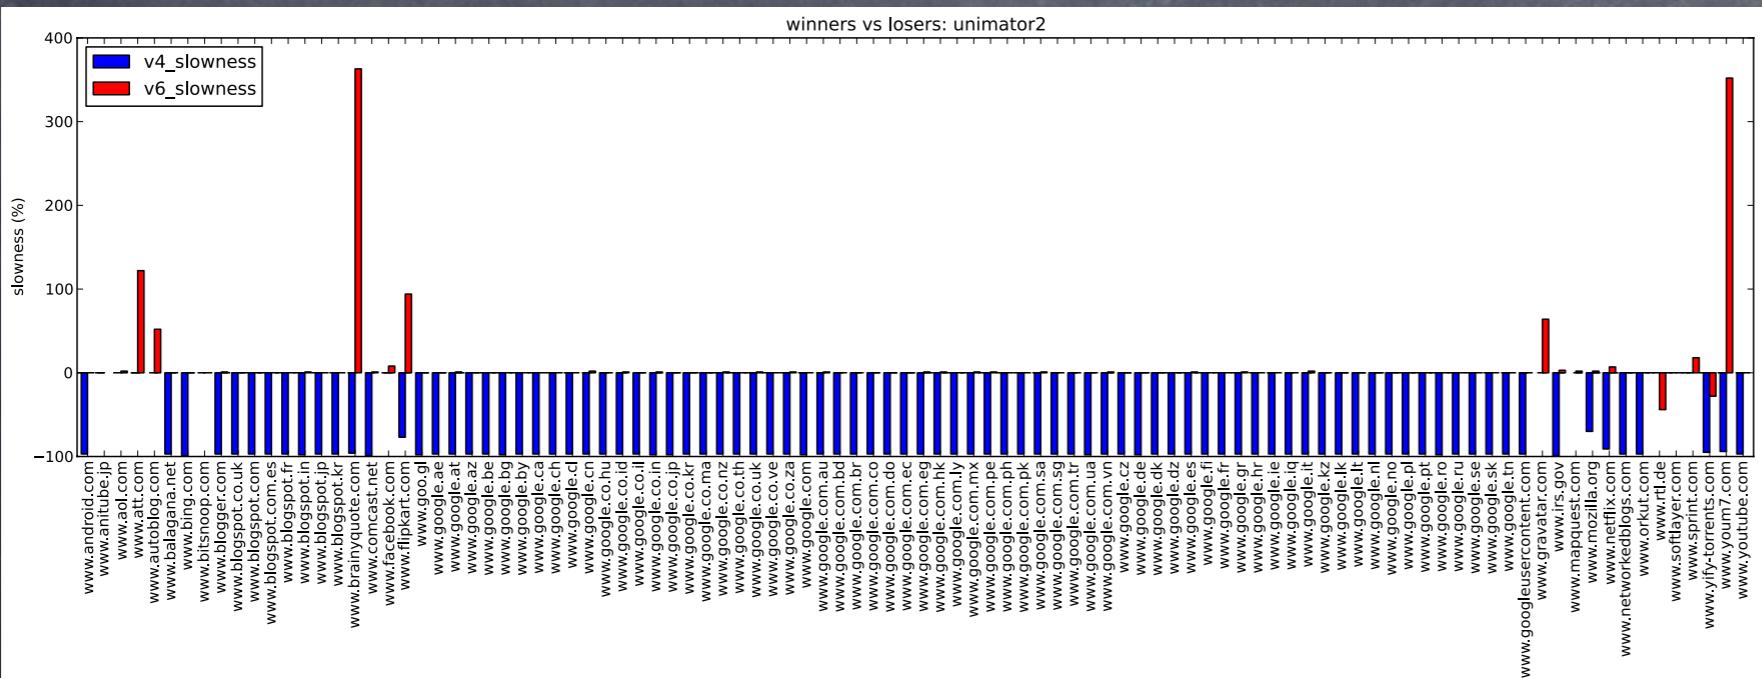
how slow is a happy eyeballed  
winner to that of a loser?

# winner slowness to loser

## Native IPv6 [Bremen]



## Native IPv6 [Braunschweig]



# conclusion

- ⦿ higher connection times and variations over IPv6
- ⦿ will never use Teredo IPv6 unless IPv4 connectivity is broken
- ⦿ 300ms advantage leaves 1% chance to prefer IPv4 (even though faster)
- ⦿ IPv6 happy eyeballed winner is rarely faster than IPv4 route
  
- ⦿ future work:
  - ⦿ provision happy on a large-scale measurement platform
  - ⦿ cluster web services using announced BGP prefixes or path similarity